

Micro-Architectural Attacks

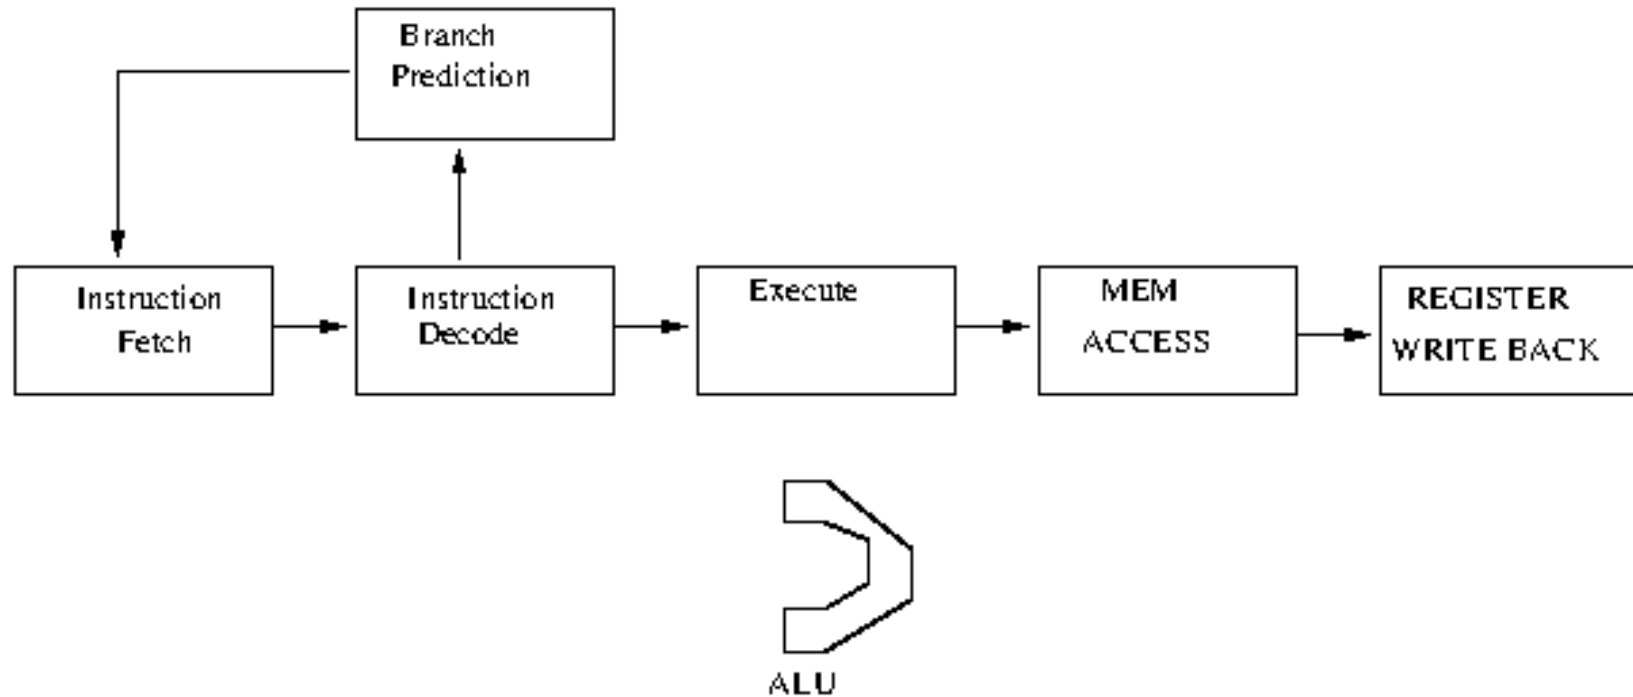
Master SETI

Sumanta Chaudhuri

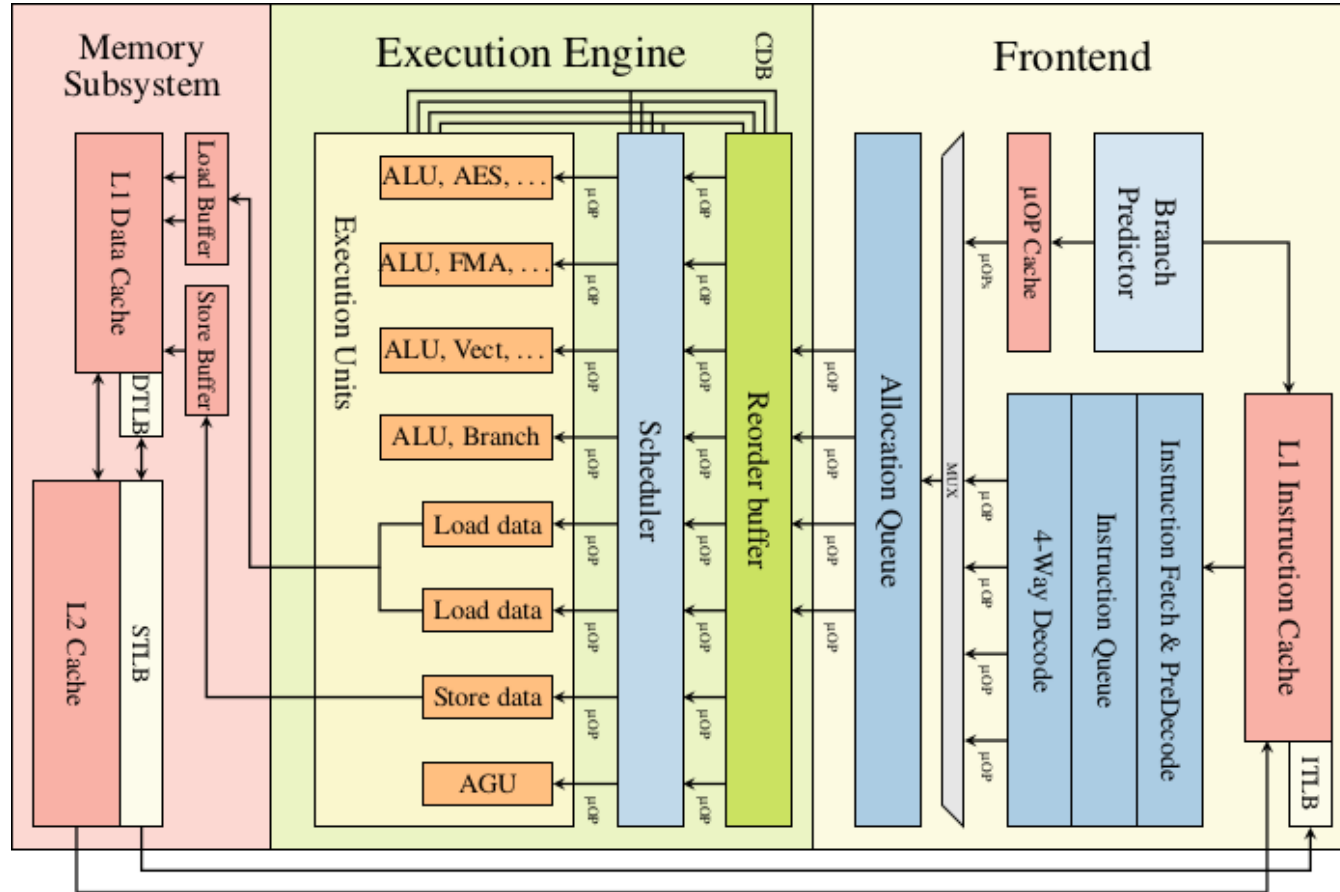
In-Order Processors

- Instruction Fetch (IF)
- Instruction Decode (ID)
- Execute (EX)
- Memory Access (MEM)
- Register Write Back (WB)

Basic gem5 In-Order Processor



Out of Order Processors



Out of Order Processors

- Multiple Instructions are fetched in parallel.
- Execute Instructions that are ready (i.e. data available)
- Instructions are committed in-order using the reorder buffer

Out of Order Processors

- Hides Latency (Like Cache, & Multiple threads)
- Much More complex
- Security Hazards (As we will see later)
- Can not be done in compiler as compiler does not have runtime data.

Out of Order Processors

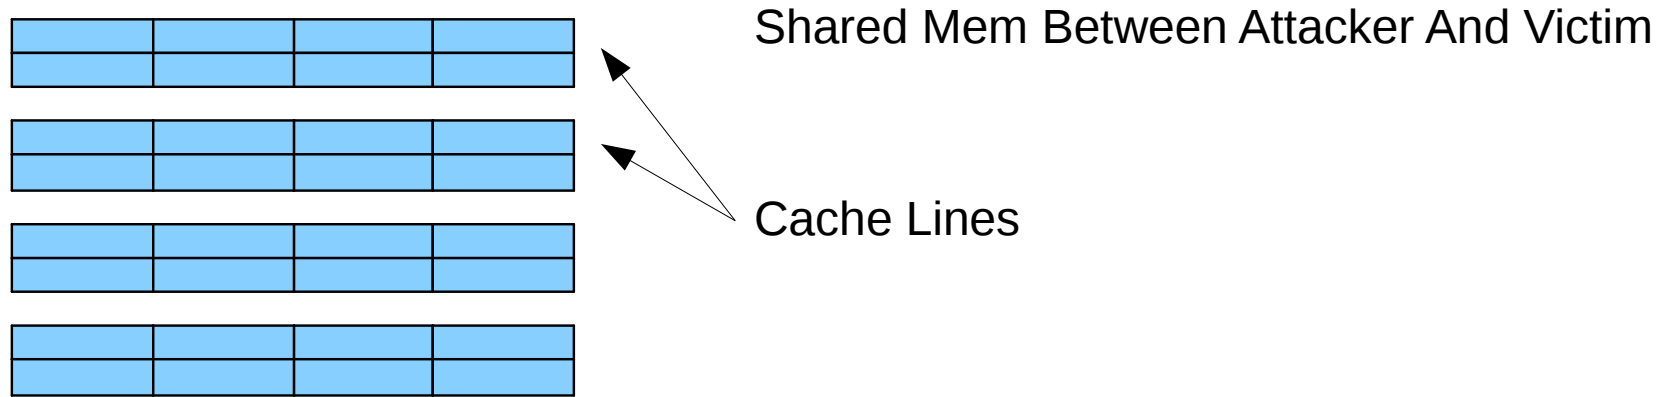
- Much longer look-ahead than in-order processors.
(e.g 150-200 instructions , Intel)
- More complex branch predictor
 - Need to predict multiple branches.
 - Several Strategies
 - Static
 - Dynamic (based on the acces pattern)
 - Misprediction Penalty: In case of misprediction all instruction of that branch have to be flushed.

Micro-Architectural Attacks

Side Channels


- Side Channels: Unintended Information leakage
 - e.g power consumption, EM emission
- Cache hit/miss Side Channels
 - The difference in hit and miss latency communicates the presence of a certain data in the shared cache.
- Speculative Execution Side Channels
 - Information leaked through speculative execution (e.g branch prediction)

Flush+Reload



Attacker **flushes** the whole array from the cache
`#include <intrin.h>`
`__mm_cflflush(array);`

Flush+Reload




Shared Mem Between Attacker And Victim

Victim acces : bring line into cache

Victim Acceses the shared cache

Flush+Reload



Shared Mem Between Attacker And Victim

Reload access time

Attacker **Re-accesses** the shared cache, Low access time due to cache hit; Measures access time
time1 = __rdtscp(array)

Figure 1 illustrates five scenarios (A-E) of attacker and victim activity. The figure shows five rows, each representing a scenario. Each row has two horizontal bars: one for the Attacker and one for the Victim. The Attacker bar is divided into segments of red (Reload), blue (Flush), and grey (Wait). The Victim bar is divided into segments of green (Access) and white (Something else). Scenario (A) shows a single reload followed by a flush. Scenario (B) shows a single reload followed by a flush. Scenario (C) shows a single reload followed by a flush. Scenario (D) shows a single reload followed by a flush. Scenario (E) shows a single reload followed by a flush.

Legend:

- Attacker:
 - Flush (Blue)
 - Wait (Grey)
 - Reload (Red)
- Victim:
 - Access (Green)
 - Something else (White)

Flush+Reload Victim




```
1 function exponent(b, e, m)
2   begin
3     x ← 1
4     for i ← 1 downto e do
5       x ← x2
6       x ← x mod m
7     if (e = 1) then
8       x ← xb
9       x ← x mod m
10    end if
11  done
12  return x
13 end
```

Flush+Reload Victim

- The instruction for RSA is in shared memory.
- The Bit '1' and bit '0' of exponent can be determined from instructions executed.
- Let's do flush+reload on the instruction cache.
 - Square+Modulo, Multiply+modulo : bit '1'
 - Square+Modulo : bit '0'


Prime+Probe

- Attacker fills the Cache with an array
-
- Tries to use the same cache lines as the victim
-
- Victim does not access the cache.
-
- Attacker Probes (reads and measure timing for his own array)
-
-
- Guesses that Victim has not accessed the target memory location




Prime+Probe

- Attacker fills the Cache with an array. (Prime)
- Tries to use the same cache lines as the victim
- Victim access the cache.
- Attacker Probes (reads and measure timing for his own array)
- Guesses that Victim has accessed the target memory location because it has evicted his own array cache line.




Prime+Probe



How to calculate the eviction sets ?
What is the offset for a cache line size of 64 bytes ?

Attack Scenario



Other Attack Scenarios can be Javascript running on browser

CounterMeasures

- Prohibit `clflush()` instruction for users.
- Prevent Sharing pages between victim and attacker code.
 - Page sharing is good for memory use.
- Detect cache attack behaviour at runtime.
- Write code to avoid cache side channel dependencies


Spectre

```
if (x < array1_size)
```

```
    y = array2[array1[x] * 4096];
```

- Speculative execution of branch even when $x > \text{array1_size}$.
- Train the branch predictor for some iterations. Force it to mispredict.

Spectre



Spectre

```
if (x < array1_size)
    y = array2[array1[x] * 4096];
```

To attack

- victim_address=array1+x
- So $\rightarrow x = \text{victim_address} - \text{array1}$
- The array2 index accessed is the value stored in victim_address.

Spectre

```
if (x < array1_size)
```

```
    y = array2[array1[x] * 4096];
```

To attack

- Find out the **array2 index accessed** with **Flush+Reload**
- Why do we need to multiply by a stride of 64 ?

Spectre Mitigations


- All Out-of-Order Processors are affected by spectre.
- However it is harder to exploit. Need to find code pattern in the victim:

```
if (x < array1_size)
```

```
    y = array2[array1[x] * 4096];
```


MeltDown

- Reading Kernel Memory from User Space



MeltDown

Exception
Handler



A rectangular box representing a processor state. It contains a sequence of five items: `<instr>`, `<instr>`, `<instr>`, `<instr>`, and `<Exception>`. A horizontal red line is drawn across the box, separating the first four instructions from the exception.

Executed

`<instr>`
`<instr>`
`<instr>`

Out of order.
Execution.

Meltdown

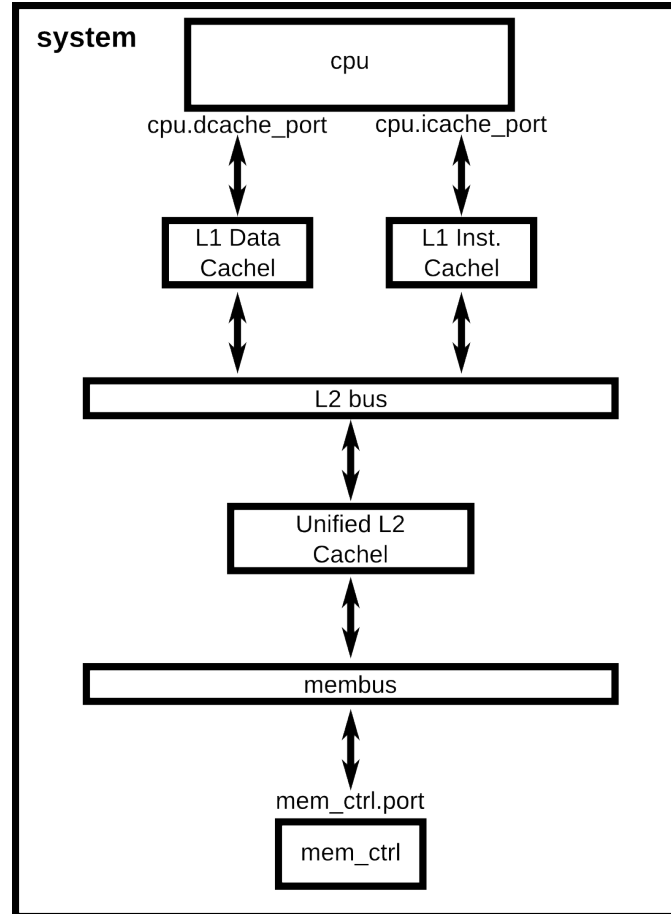
1. `raise_exception();`
2. `// the line below is never reached`
3. `access(probe_array[data * 4096]);`

Spill over to the Kernel memory space.
Find the value through Flush+Reload.

Meltdown Mitigations

- KAISER Patch: User space does not have access to kernel memory.
- KASLR (Address space layout randomization): Makes the attack difficult.

TP: CONFIG



TP: STEP 1

- Clone the repository https://github.com/amusant/micro_archi_attacks
- `$source env.sh` → sets up environment variables.
- Go to directory `hit_miss`; look into code `hit_miss.c`
- Run make to compile the code in `hit_miss` directory
- Runs \$make launch to launch simulation.
-
- We use the gem5 simulator to simulate a basic system with x86 processor and two levels of cache.
- Understand the code used for
 - Flush
 - Acces
 - Reload
- **By changing the acces pattern do you see any difference in the output ?**
- **What is the role of STRIDE, does the code still work after changing STRIDE ?**

TP: STEP 2

- Go to directory `flush_reload`; look into code `flush_reload.c`
-
- The function `victim(k)` does the following:
 - It accesses the **`array[secret[desknumber][i]*STRIDE]`**
 - Where the secret is a 16 character secret key.
 - **`secret[desknumber]="XXXXXXXXXXXXXXXXXX"`**
 - Your goal is to find the 16 characters of the secret value.
 - The secret value changes with desk number.
- Run `make` to compile the code in `flush_reload` directory
- Runs `$make launch` to launch simulation.
- **Inspire yourself from the `hit_miss` code.**

TP: STEP 3

- Download the Spectre Example here
- <https://gist.github.com/ErikAugust/724d4a969fb2c6ae1bbd7b2a9e3d4bb6>
-
- Read and Understand the code.
- Compile the code
 - **\$gcc spectre.c**
 - Launch the experiment
 - **\$ gem5.opt ../configs/two_level.py ./a.out**
 - Does it work ?
- Change line 99 in **../configs/two_level.py**
 - from `DerivO3CPU()` to `TimingSimpleCPU()`
 - - relaunch simulation
 - - Does it work ?